

# Speeding Up Packet I/O in Virtual Machines

Luigi Rizzo  
Università di Pisa, Italy  
rizzo@iet.unipi.it

Giuseppe Lettieri  
Università di Pisa, Italy  
lettieri@iet.unipi.it

Vincenzo Maffione  
Università di Pisa, Italy  
v.maffione@gmail.com

## ABSTRACT

Most of the work on VM network performance has focused so far on bulk TCP traffic, which covers classical applications of virtualization. Completely new “paravirtualized devices” (Xenfront, VIRTIO, vmxnet) have been designed and implemented to improve network throughput.

We expect virtualization to become widely used also for different workloads: packet switching devices and middleboxes, Software Defined Networks, etc.. These applications involve very high packet rates that are problematic not only for the hypervisor (which emulates network interfaces) but also for the host itself (which switches packets between guests and physical NICs).

In this paper we provide three main results. First, we demonstrate how rates of millions of packets per second can be achieved even within VMs, with limited but targeted modifications on device drivers, hypervisors and the host’s virtual switch. Secondly, we show that emulation of conventional NICs (e.g., Intel e1000) is perfectly capable of achieving such packet rates, without requiring completely different device models. Finally, we provide sets of modifications suitable for different use cases (acting only on the guest, or only on the host, or on both) which can improve the network throughput of a VM by 20 times or more.

These results are important because they enable a new set of applications within virtual machines. In particular, we achieve guest-to-guest UDP speeds of over 1 Mpps with short frames (and 6 Gbit/s with 1500-byte frames) using a conventional e1000 device, and socket-based sender/receivers. This matches the speed of the OS on bare metal. Furthermore, we reach over 5 Mpps when guests use the netmap API.

Our work requires only small changes to device drivers (about 100 lines, both for FreeBSD and Linux version of e1000), similarly small modifications to the hypervisor (we have a QEMU prototype available) and the use of the VALE switch as a network backend. Relevant changes are being incorporated and/or distributed as external patches for FreeBSD, QEMU and Linux.

## Categories and Subject Descriptors

D.4.4 [Operating Systems]: Communications Management  
Network Communication

## General Terms

Design, Experimentation, Performance

## Keywords

Software Switches; Virtual Machines; netmap

## 1. INTRODUCTION

Virtualization is a technology in heavy demand to implement server consolidation, improve service availability, and make efficient use of the many cores present in today’s CPUs. Of course, users want to exploit the features offered by this new platform without losing too much (or possibly, anything) of the performance achievable on traditional, dedicated hardware (*bare metal*).

Over time, ingenious software solutions [5], and later hardware support [13, 3], have mostly filled the gap for CPU performance. Likewise, performance for storage peripherals and bulk network traffic is now comparable between VMs and bare metal, especially when I/O can be coerced to use large blocks (e.g., through TSO/RSC) and limited transaction rates (e.g., say less than 50 K trans/s).

However a class of applications, made relevant by the rise of Software Defined Networking (SDN), still struggles under virtualization. Software routers, switches, firewalls and other middleboxes, need to deal with very high packet rates (millions per second) that are not amenable to reduction through the usual Network Interface Card (NIC) offloading techniques. The “direct mapping” of portions of virtualization-aware NICs to individual VMs can provide some relief, but it has scalability and flexibility constraints.

We then decided to explore solutions to let VMs deal with millions of packets per second without requiring special hardware, or imposing massive changes to OSes or hypervisors. In this paper we discuss the general problem of network performance in virtual machines, identifying the main causes of performance loss compared to bare metal, and design and experiment with a comprehensive set of mechanisms to fill the performance gap.

**Our contribution:** in detail, we i) emulate interrupt moderation, ii) implement “Send Combining”, a driver-based form of batching and interrupt moderation; iii) introduce an extremely simple but very effective paravirtualized extension for the e1000 devices (or other NICs), providing the same performance of VIRTIO and alikes with almost no extra complexity; iv) adapt the hypervisor to our high speed VALE [20] backend, and v) characterize the behaviour of *device polling* under virtualization.

Some of the mechanisms we propose help immensely, especially within packet processing machines (software routers, IDS, monitors ...). Especially, the fact that we provide solutions that apply only to the guest, only to the host, or to




Figure 1: In our virtualized execution environment a virtual machine uses one VCPU thread per CPU, and one or more I/O threads to support asynchronous operation. The hypervisor (VMM) has one component that runs in userspace (QEMU) and one kernel module (KVM). The virtual switch also runs within the kernel.

both, makes them applicable also in presence of constraints (e.g., legacy guest software that cannot be modified; or proprietary VMMs).

In our experiments with QEMU-KVM and `e1000` we reached a VM-to-VM rate of almost 5 Mpps with short packets, 25 Gbit/s with 1500-byte frames, and even higher speeds between a VM and the host. These large speed improvements have been achieved with a very small amount of code, and our approach can be easily applied to other OSes and virtualization platforms. We are pushing the relevant changes to QEMU, FreeBSD and Linux.

In the rest of this paper, Section 2 introduces the necessary background and terminology on virtualization and discusses related work. Section 3 describes in detail the components of our proposal, whereas Section 4 presents experimental results, and also discusses the limitations of our work.

## 2. BACKGROUND AND RELATED WORK

In our (rather standard) virtualization model (Figure 1), Virtual Machines (VMs) run on a Host which manages hardware resources with the help of a component on the Host called *hypervisor* or Virtual Machine Monitor (VMM, for brevity). Each VM has a number of Virtual CPUs (VCPU, typically implemented as threads in the host), and also runs additional I/O threads to emulate and access peripherals. The VMM (typically implemented partly in the kernel and partly in user space) controls the execution of the VCPUs, and communicates with the I/O threads.

The way virtual CPUs are emulated depends on the features of the emulated CPU and of the host. The x86 architecture does not lend itself to the *trap and emulate* implementation of Virtualization [1], so historical VMMs (VMware, QEMU) relied for the most part on binary translation for “safe” instructions, and calls to emulation code




Figure 2: Different ways to emulate IO register accesses. They can be taken care of directly in the kernel component of the hypervisor (case “a”), or in the host component (case “b”), or handed off to the I/O thread (case “c”) in case of long activities.

for others. A recent paper [5], long but very instructive, shows how the x86 architecture was virtualized without special support for virtualization on the host CPU. The evolution of these techniques is documented in [1]. A slowdown of 2..10 times can be expected for typical code sequences that run in the virtual machine; slightly better performance can be achieved if kernel support is available to intercept memory accesses to invalid locations.

Modern CPUs provide hardware support for virtualization (Intel VT-X, AMD V) [13, 3], so that most of the code for the guest OS is run directly on the host CPU operating in “VM” mode. In practice, the kernel side of a VMM enters VM mode through a system call (typically an `ioctl(... VMSTART ...)`), which starts executing the guest code within the VCPU thread, and returns to host mode as described below.

### 2.1 Device emulation

Emulation of I/O devices [25] generally interprets accesses to I/O registers and replicates the behaviour of the corresponding hardware. The VMM component reproducing the emulated device is called *frontend*. Data from/to the frontend are in turn passed to a component called *backend* which communicates with a physical device of the same type: a network interface or switch port, a disk device, USB port, etc.

Access to peripherals from the guest OS, in the form of IO or MMIO instructions, causes a context switch (“VM exit”) that returns the CPU to “host” mode. VM exits often occur also when delivering interrupts to a VM. On modern hardware, the cost of a VM exit/VM enter pair and IO emulation is 3..10  $\mu$ s, compared to the 100-200 ns for IO instructions on bare metal.

The detour into host mode is used by the VCPU thread to interact with the frontend and emulate the actions that the real peripheral would perform on that specific register access. In some cases these are trivial (such as reading or writing a memory word); other times, a register access may involve copying one or more blocks of data, e.g. network frames, between the emulated device and the actual peripheral on the host.

This emulation can be done in different ways, as illustrated in Figure 2. For frequently accessed peripherals that

are a standard part of an architecture, such as interrupt control registers, the kernel side of the VMM can do the emulation directly; this is represented by case “a” in Figure 2).

In other cases, the task can be performed in the user side of the VMM (see case “b”), which gives greater flexibility but also adds cost to the operation. As an example in one of our systems (Intel i5-2450M @ 2.5 GHz, memory @ 1333 MHz) a register access requires  $1.2\mu\text{s}$  when executed by the kernel component of the hypervisor (KVM), and  $3.7\mu\text{s}$  when dispatched to the userspace component.

In the worst case, the VCPU thread may actually even be descheduled, or block waiting for I/O. This can be problematic as I/O accesses often occur under some kind of locking, so changing the execution time from 100 ns to potentially much longer times may significantly reduce the ability of other VCPU threads to perform useful work. It follows that potentially long operations are better served by running them within an independent thread (I/O thread), to which the VCPU hands off the request (case “c”). I/O threads are also used to handle events (e.g., I/O completions, timeouts) generated from the host OS.

This description illustrates the fundamental performance issues and tradeoffs in using emulated peripherals within a VM. While regular code runs at full speed, IO instructions consume thousands of clock cycles, and can block other VCPUs for very long times. Also, the latency in performing the actual I/O operation is badly affected by the thread handoffs that are necessary to return quickly from a VM exit.

Hence, a significant challenge in VM environments is the reduction or elimination of VM exits, which are especially frequent when dealing with high packet rate network traffic. Our work is actually providing a number of solutions in this space.

## 2.2 Reducing VM exits

It is not uncommon that instructions leading to VM exits are close to each other – as an example, within device drivers where multiple registers are accessed at once. A recent proposal [2] shows how short sequences of code involving multiple I/O instructions can be profitably run in interpreted mode to save some VM exits. Implementing this technique is non trivial, but it can be done completely within the VMM without any change in the guest OS.

VM exits are not needed for hardware that is directly accessible to the guest VM: we can give it full control of a physical device (or part of it, if the device supports “virtual functions”) as long as we can make sure that the VM cannot fiddle with other reserved parts of the hardware. IOMMU [4] and related technologies (generally referred to as *PCI passthrough*) comes to our help here, as they provide the necessary protections as well as translation between the guest’s physical address space and the host physical address space. Again, this approach can be almost completely transparent for the guest OS, and relies only on support in the hardware and the VMM.

### 2.2.1 Paravirtualized devices

Another popular approach to reduce VM exits is to design new device models more amenable to emulation. This approach, called “paravirtualization”, has produced several NIC models (vmxnet [26], VIRTIO [21], xenfront [6]), in turn requiring custom device drivers in the guest OS. Syn-

chronization between the guest and the VMM uses a shared memory block, as described in Section 3.5, accessed in a way that limits the number of interrupts and VM exits. One contribution of this paper is to show that paravirtualization can be introduced with minimal extensions into existing NICs.

### 2.2.2 Interrupt handling

Interrupts to the guest frequently cause exits, either directly or indirectly (because they trigger accesses to registers of interrupt controllers and devices). Mechanisms to reduce interrupt rates, as those shown in Section 3.1 and following, help significantly.

Likewise, it is possible for interrupts to be delivered directly to the guest without causing VM exits. As an example, ELI [8] removes interrupt-related VM exits on direct-access peripherals by swapping the role of host and guest: the system is programmed so that all interrupts are sent to the guest, which reflects back those meant for the host.

## 2.3 High speed networking

Handling 10 Gbit/s or faster interfaces is challenging even on bare metal. Packet rates can be reduced using large frames, or NIC/OS support for segmentation and reassembly (named TSO/GSO and RSC/LRO, respectively). These solutions do not help for packet processing devices (software routers, switches, firewalls), which are not sourcing or terminating connections, hence must cope with true line speed in terms of packet rates.

Only recently we have seen software solutions that can achieve full line rate at 10 Gbit/s on bare metal [18, 7, 12]. In the VM world, apart from the trivial case of directly mapped peripherals, already discussed [8], the problem has not seen many contributions in the literature so far. Among the most relevant:

Measurements presented in [23] show that packet capture performance in VMs is significantly slower than on native hardware, but the study does not include the recent techniques mentioned above, nor proposes solutions. Interrupt coalescing and Virtual Receive Side Scaling have been studied in [10] within Xen; the system used in that paper is limited to about 100 Kpps per core, and the solutions proposed impose a heavy latency/throughput tradeoff and burn massive amounts of resources to scale performance. ELVIS [9] addresses the reduction of VM exits in guest-host notifications: a core on the host monitors notifications posted by the guest(s) using shared memory, whereas inter-processor interrupts are used in the other direction, delivered directly to the guest as in the ELI case.

## 2.4 State of the art in VM networking

On bare metal and suitably fast NICs, clients using a socket API are generally able to reach 1 Mpps per core, peaking at 2.4 Mpps per system due to driver and OS limitations. Recent OS-bypass techniques [18, 7, 14, 24] can reach much higher rates, and are generally I/O bound, easily hitting line rate (up to 14.88 Mpps on a 10 Gbit/s interfaces) unless limited by NIC or PCIe bus constraints.

Within a VM, at least when emulating regular NICs (e.g., the popular Intel’s *e1000*), common VMMs reach rates of about 100 Kpps on the transmit side, and marginally higher rates on the receive side (see Figure 4, 5 and 10). Paravirtualized devices (VIRTIO etc.) can help significantly, but support for them is neither as good nor as widespread as

it is for popular hardware NICs. This is especially true in the presence of old/legacy systems that cannot be decommissioned: they are often run within a VM due to lack of suitable hardware. So there is still a strong motivation for efficient NIC emulation, which motivates some of the solutions that we propose, requiring none or minimal modifications to the guest OS — much less than adding an entirely new device driver.

Note that our work goes beyond a simple replacement of paravirtualized devices. We address packet rates beyond those supported by the usual virtual switches, so we needed to identify and fix performance bottlenecks in multiple places of the architecture, from the guest operating system to the virtual switch itself.

### 3. OUR WORK

Taking QEMU-KVM as a prototype platform, we have investigated how to improve the network performance of guests, using both regular (e.g., Intel e1000) and paravirtualized network devices. Given the huge number of dimensions in the problem space, we initially focus the presentation on the cases in dire need of improvement: guest to guest communication, high packet rates, non-paravirtualized devices. Other configurations will be discussed and evaluated in Section 4.

For the host OS we used Linux 3.8, with KVM. For the guest OS we used both Linux 3.8 and FreeBSD HEAD. The differences in device driver and network stack architectures helped pointing out different phenomena and also the peculiarities in running in a VMs rather than on bare metal.

We should also keep in mind the constraints that may limit the solution space. Sometimes the guest OS cannot be modified or extended, hence we can only operate on the host side; in other cases the hypervisor is instead proprietary, hence we should look for guest-only solutions. And of course, in absence of such constraints, we should look for performance improvement techniques across the board. We believe that the set of solutions presented in this paper covers nicely the various situations, and can be easily applied to other systems/platforms than the ones we used for our experiments.

#### 3.1 Interrupt moderation on emulated NICs

Interrupt moderation [15] is a widely used technique to reduce interrupt load, and it is normally available on modern NICs and supported by OSes. It struck our attention that the feature is rarely implemented by VMMs: QEMU and VMware Player do not emulate moderation registers; VirtualBox supports them only in very recent versions.

This was surprising since interrupts storms from NICs at high packet rates are a significant source of uncontrolled system load, leading to receiver livelock and severe reductions of the transmit rate.

One reason for the omission may be that modern Operating Systems employ various interrupt avoidance techniques in software (e.g., the NAPI architecture in Linux [22], device polling in FreeBSD [17]; lazy buffer recovery on the transmit side, see Sec. 3.3). However, except for [17], such techniques do not provide guarantees on the maximum interrupt rate, so interrupt moderation is still useful to avoid livelock, though not as much as for older OSes and drivers that do not exploit interrupt avoidance techniques. Also, having interrupt moderation in the device enables the use of

Send Combining (discussed next in Sec. 3.2) on the transmit side which has significant benefits in a VM.

Our first modification was to implement the interrupt moderation registers in QEMU. The (small) change of the code affects only the hypervisor, and can work even if guest OS or the original NIC do not support moderation, as we can enforce the feature on the emulated device. Hence this makes a good solution for old or not modifiable guests.

##### 3.1.1 Effects on transmission

On the transmit side, moderation more than doubles the transmit rate for UDP, especially with a single VCPU (on Linux, we go from 29 to 82 Kpps with 64-byte frames, 28 to 70 Kpps with 1500-byte frames). The reason is that on the VM, in absence of moderation, a transmit request (triggered by a write to one of the NIC registers, in turn causing a VM exit) is instantly followed by an interrupt, that causes a burst of VM exits in the interrupt service routine. As a consequence, each interrupt serves exactly one packet. With moderation, one exit per packet is still used in the transmit routine, but the cost of the interrupt routine is now amortized over a larger set of packets.

The effect of moderation is more limited with 2 VCPUs, because in this case the interrupt service routine can run in parallel with the transmit requests, and this gives some chances to serve more than one packet per interrupt. Depending on the OS, we may still see some improvement, but more limited, e.g., Linux goes from 40 to 82 Kpps.

Note that adding a second VCPU increases the base throughput from 29 to 40 Kpps, but it has little or no effect when moderation is used (we have observed about 82 Kpps in both cases). This can be explained as follows: the effect of the second VCPU is to offload some interrupt processing from the first VCPU thread, which is the bottleneck in these benchmarks; but the second VCPU is only activated by the NIC generated interrupts, and therefore it has much less opportunities to kick in if interrupt moderation is used. This can be confirmed by looking at the CPU utilization of the second VCPU thread in the two scenarios: 56% in the base case, but only 13% when moderation is turned on. This much smaller effect is likely counterbalanced by increased overheads.

##### 3.1.2 Effects on reception

On the receive side, the effect of moderation depends on the arrival pattern of traffic. In some cases, traffic received by the VM already comes in batches (with one interrupt per batch) because this is how it is passed on by the incoming (physical) interface, virtual switch and/or the backend. Yet, even in these cases, the use of interrupt moderation can help increase the batch size and further amortize the cost of interrupt processing, and increase the rate at which receive livelock may appear (see Section 3.4).

### 3.2 Send Combining

Device drivers typically notify the NIC immediately when new packets are ready for transmission; this is relatively inexpensive on bare metal, but causes a VM exit in a VM.

To reduce the number of such exits we used an idea similar to what in [25, Sec.3.3] is called “Send Combining” (SC): since the driver knows about pending TX interrupts, it can defer transmission requests until the arrival of the interrupt. At that point, all pending packets are flushed (with a single



write on one of the NIC’s registers), so there is only one VM exit per batch.

SC is a trivial (about 25 lines of code in our case), *guest only* modification of the transmit path. However it has no effect if TX interrupts are instantaneous, so it is mostly useful when coupled with moderation, or with device polling, see Section 3.3.2, because it reduces the system load from one VM exit per packet to one VM exit per batch.

On the negative side, send combining may delay transmissions by up to the moderation/polling delay (20..100  $\mu$ s for moderation, 250..1000  $\mu$ s for polling), similarly to what happens on the receive path with these techniques. The inconvenience can be reduced by tuning SC (and interrupt moderation) to kick in only above a certain packet rate, and by choosing appropriate tradeoffs between delay and throughput.

The combination SC plus moderation gives a fantastic boost on the transmit path: on Linux, we went from 40 to 354 Kpps with 2 VCPUs, and from 30 to 221 Kpps with 1 VCPU. The reason for these high speedups is the massive reduction of VM exits (by far the most expensive operation), which are now well below one per packet. We are now approaching rates where the virtual switch and the data path within the VMM may become a performance bottleneck.

As a final note in this section, we should mention that unlike our version, the original Send Combining in [25] was implemented entirely in the VMM, because it ran I/O instructions in “binary translation” mode. When using CPU-based virtualization support we do not have this luxury, and we cannot intercept I/O requests before they cause a VM exit.

### 3.3 Reducing interrupts


Interrupt moderation in the emulated device is not the only way to reduce interrupts. There are other techniques that can be implemented in the guest OS, and lead to similar or better performance improvements. The following ones are not new contributions of this work, but it is important to understand how they relate to virtualization. Especially, the benefits of polling within VMs has not been documented so far.

#### 3.3.1 Lazy TX completion

As long as there are resources available, it is not necessary to interrupt on each transmitted packet to recover the buffer. *Lazy TX completion* recovers completed transmission in batches (we can think of it as moderation based on packet counts rather than time), or opportunistically during subsequent transmissions. This technique is implemented in some device drivers, and also in the netmap [18] framework. Note that it can almost completely remove transmit interrupts (they are still necessary when the transmit queues are full, as a notification that new transmit resources are available), so it is not compatible with send-combining, which instead relies on interrupts to flush pending transmissions.

#### 3.3.2 FreeBSD’s device polling

The FreeBSD’s device polling framework [17] completely disables NIC interrupts, and runs the equivalent of the interrupt service routines at every timer tick (0.2 .. 1 ms on most systems) and optionally in the system’s idle thread. The framework was designed for a totally different purpose (enabling moderation for devices that do not support it, and



**Figure 3: UDP receive throughput for different input traffic. Only Linux-VIRTIO (without VHOST) and FreeBSD polling are able to sustain their peak rate for input rates up to 20 Mpps (the graph is truncated to 6 Mpps for clarity).**

preventing livelock under high network load), but turns out to be extremely effective when used within a virtual machine.

In fact, by removing interrupts, polling also removes the many VM exits that are generated to access the interrupt controller registers, so it is even more effective than interrupt moderation. Coupled with send combining, device polling also achieves great rates on the transmit path. As an example, just enabling polling on FreeBSD, with proper parameters, supports receive rates of up 890 Kpps, and to about 450 Kpps on the transmit side.

### 3.4 Receive livelock

Receive livelock [15] refers to a significant throughput decrease as the input rate grows. It generally arises when incoming traffic is not processed to completion in a single stage, and the initial processing stage (e.g., interrupt handling) is given priority over others.

Livelock is a serious problem in high speed network environments. Interrupt moderation, coupled with modern, powerful CPUs, is often sufficient to alleviate the problem.

Within VMs, the problem is aggravated by two factors: the much higher cost of I/O operations, and the fact that emulated hardware can be subject to much higher input rates than the physical devices they emulate. As an example, the `e1000` driver normally manages 1 Gbit/s cards, resulting in at most 1.488 Mpps. However in a VM the driver can receive much higher packet rates, e.g. the VALE[20] backend presented in Section 3.6 is able to drive it with up to 20 Mpps and 70 Gbit/s. This defeats all livelock prevention mechanisms that rely on abundance of CPU resources.

To illustrate the phenomenon, Figure 3 shows the receive throughput for an UDP socket under variable input load (with input coming in bursts of 100 packets), with different combinations of operating system (FreeBSD or Linux), device driver (`e1000` or VIRTIO), operating mode (NAPI, paravirtualized driver as described in Section 3.5, FreeBSD polling [17]), and network backend (TAP and VALE, described in Section 3.6).

For most configurations (including several not shown) the receive rate drops to very low values, or even to 0, as the input rate grows. In our experiments, only two solutions resisted to extreme input rates without incurring in livelock: the FreeBSD polling framework and Linux VIRTIO (without VHOST).

The robustness of FreeBSD polling comes from two reasons: first, the maximum interrupt rate is limited to a relatively low rate (1.20 K per second), which does not pose problems even with the expensive VM exits; second, a control mechanism dynamically adjusts the work performed in the interrupt routines to guarantee a configurable amount of CPU time to other processes. This makes sure that, even under extreme input loads, the desired fraction of CPU capacity is still available to perform work and complete processing of the input.

The robustness of VIRTIO-based configurations is instead accidental, and derives from the inability of the VIRTIO QEMU frontend (which in the non-VHOST case runs in the userspace emulator) to deliver more than 700Kpps (in our experiment) to the VM. As a consequence, in this configuration we do not see livelock only because the guest is never running out of CPU. When VIRTIO is accelerated with VHOST, the VIRTIO frontend (which is now in the host kernel) is no more the bottleneck, and we do see the livelock for high rates.

### 3.5 Paravirtualizing real NICs

Paravirtualized devices (VIRTIO, vmxnet, xenfront) generally reduce the number of VM exits by establishing a shared memory region (we call it *Communication Status Block* or CSB) through which the guest and the VMM can exchange I/O requests without VM exits.

The mechanism works as follows: on the first packet after an idle interval, the traffic source (either the guest or the I/O thread) issues an initial notifications (called *kick*) to wake up the other peer. Kicks from the guest are writes to a NIC register, and do cause a VM exit; kicks from the host (typically issued by the I/O thread) are interrupts, and cause VM exits too. The entity that receives a kick polls the CSB to look for new work (new packets available), and posts indications to the CSB as they are processed. The polling continues until there is no more pending work, at which point the entity reports in the CSB that is about to go to sleep, and actually does so when no more work is available for a while.

Same as other mechanisms discussed earlier, paravirtualization reduces the number of VM exits to one (or a few) per batch of packets, and becomes more effective as the load increases, a nice scaling property. The mechanisms required to implement paravirtualization do not differ much from what is already available in modern NICs; in fact:

- the hardware already reports packet receptions and transmit completions through status bits in the descriptor ring, which is accessible through shared memory without VM exits;
- the register writes that start transmissions, and the receive interrupts are perfectly equivalent to a “kick” in VIRTIO terminology (in fact these two mechanisms are used to implement VIRTIO as a PCI device, which is the implementation used in the VMMs).

The only missing features to support paravirtualization are copies of the registers used to indicate new transmissions and new buffers for the receive ring, which the guest can update without VM exits.

Hence all it takes to build a full paravirtualized device is a small region of memory to implement the CSB (which includes the two registers mentioned above, plus some extra control information), and the code to exchange information through it.

We have implemented paravirtualization support for the e1000 and r8169 devices and drivers by adding two PCI registers to the emulated NIC to point to the mapped CSB. The new capability is advertised through the PCI subdevice ID, and the device driver must explicitly enable the feature for the emulator frontend to use it.

Overall, this modification required about 100 lines of code in the guest device driver, and approximately the same amount in the frontend. The same CSB structure is used for both NICs.

This is immensely simpler than writing an entirely new device driver (about 3000 lines for VIRTIO-net) and frontend (about 1500 lines). In terms of performance, our paravirtualized NICs perform as well as VIRTIO. As such, this is a more viable approach to extend both guests and hypervisors that do not have paravirtualized devices (or have incompatible ones: it is amusing that even among paravirtualized NICs, each manufacturer has gone its own way, leading to the existence of at least three contenders — VIRTIO, vmxnet, xenfront — which replicate similar functionality).

### 3.6 VALE, a fast backend

As the experiments in the next Section will show, with the enhancements described so far have become faster than the maximum packet rates supported by the backends (sockets, tap, host bridges, ...) that interconnects virtual machines. This bottleneck used to be hardly visible due to general slowness in the guest and device emulation, but the problem clearly emerges now.

Some room for improvement still exists: as an example, the VHOST feature [11] for VIRTIO avoids going through the emulator userspace code for sending and receiving packets, and instead does the forwarding directly within the host kernel. In our experiments VHOST almost doubles the peak pps rate over TAP for the VIRTIO cases, slightly above our best result with e1000-CSB (which still does not have an equivalent optimization).

Nevertheless, our target of several millions of packets per second requires a substantially faster backend. We have then implemented a QEMU backend that uses the VALE [20] software switch, which in itself is capable of handling up to 20 Mpps with current processor technology. Attaching QEMU to VALE pointed out a number of small performance issues in the flow of packets through the hypervisor (mostly, redundant data copies and lookups of information that could be cached). Our initial implementations could “only” reach 2.5 Mpps between guest and host, but this value has been more than doubled in the current prototype.

## 4. EXPERIMENTAL RESULTS

We have anticipated some performance numbers in the previous Section, but here we provide a more comprehensive comparison of the various mechanism presented. Exploring all possible combinations of components (guest OS;

virtual CPUs; NICs; hypervisors; backends; virtual switch; load conditions) would be prohibitively time consuming, so we restrict our analysis to a (still large) number of relevant configurations.

We present now some results on the data rates we achieved between Guest and Host (GH), and between two guests (GG), with different emulators and combinations of features. Our source code containing all QEMU and guest modifications is available at [19].

### General notes

For basic UDP, TCP and latency tests we have used the popular **netperf** program. A socket-based program (**net-send** from FreeBSD) has been used to generate rate-limited UDP traffic. For very high speed tests, exceeding the data rates achievable with sockets, we have used the **pkt-gen** program part of the netmap [16] framework. **pkt-gen** accesses the NIC (or the VALE port) bypassing the network stack, acts as a sink or as a source with programmable rate, and can sustain tens of millions of packets per second.

Our host system uses an Intel i7-3930K CPU @ 3.20GHz running Linux 3.8.11; QEMU-KVM is the git-master version as of May 2013, extended with all the mechanisms described in this paper. We use **tap** or **VALE** as backends.

Our guests are FreeBSD HEAD or Linux 3.8, normally using the **e1000** driver with small extensions to implement Send Combining (SC) and Paravirtualization (CSB). The interrupt moderation delay (i.e., the minimum interval between two interrupts) is controlled by the **itr** parameter (in 256 ns units).

We also ran some tests using other hypervisors (VirtualBox, VMware Player) and/or features (VIRTIO, VHOST, TSO) to have some absolute performance references and evaluate the impact of the features we are still missing.

**Note:** the goal of this paper is to study the behaviour of the system *at high packet rates*, as those that may occur in routers, firewalls and other network middleboxes. Hence, our main focus are streams of UDP or raw packets. TCP throughput is only reported for completeness, but we did not try or want to optimize TCP parameters (buffer and window sizes, path delays, etc.) for maximum throughput but rely on system defaults, for good or bad as they are. More details on TCP performance are in Section 4.3.

### Notations and measurement strategy

In the following tables of results and related discussions, configurations are labeled according to the main parameters that have an influence on results, namely:

- the guest operating system (Linux or FreeBSD);
- the **VMM** (typically QEMU, but we ran some limited tests with VMware Player and VirtualBox);
- the **backend** (TAP or or VALE);
- whether communication is between a Guest and the Host (GH), or between two Guests VMs (GG). GH measurements are useful to evaluate separately the transmit and receive path (the host being generally a faster source/sink). Tests between two guests (GG) show the end-to-end behaviour (on the same host);
- the number of **VCPUs per guest**;

Linux on QEMU-KVM, TAP backend, Guest-Host						
1 VCPU	UDP8 Kpps	UDP-1460 Kpps Mbps	TCP Mbps	TCPw Mbps	TCPRR KTps	
TX itr=0 BASE	30	28	326	103	182	11.8
TX itr=250	83	71	827	135	249	14.2
TX itr=250 SC	216	138	1613	1309	1521	4.6
TX itr=0 CSB	353	242	2828	2620	2594	21.2
TX itr=250 CSB	360	239	2794	2590	2757	7.6
RX itr=0 BASE			3370	3327	12.3	
RX itr=100 SC			3912	3890	13.4	
RX itr=100 CSB			7179	7493		
2 VCPU	UDP8 Kpps	UDP-1460 Kpps Mbps	TCP Mbps	TCPw Mbps	TCPRR KTps	
TX itr=0 BASE	45	38	452	239	349	13.5
TX itr=250	84	54	630	329	396	13.5
TX itr=250 SC	291	136	1592	1723	1761	11.3
TX itr=0 CSB	369	190	2223	1993	1959	20.7
TX itr=250 CSB	343	172	2011	1690	1683	10.4
RX itr=0 BASE			3938	4118		
RX itr=100			1250	2894		
RX itr=100 SC			2300	5128		
RX itr=100 CSB			4400	7566		
Linux on QEMU-KVM, TAP backend, Guest-Guest						
1 VCPU	UDP8 Kpps	UDP-1460 Kpps Mbps	TCP Mbps	TCPw Mbps	TCPRR KTps	
TX = RX						
itr=0 BASE	29	28	321	416	462	7.1 *
itr=250	82	70	827	522	598	6.5
itr=250 SC	221	138	1615	1265	1661	5.1
itr=0 CSB	458	269	3147	2923	2982	11.8 *
2 VCPU	UDP8 Kpps	UDP-1460 Kpps Mbps	TCP Mbps	TCPw Mbps	TCPRR KTps	
TX = RX						
itr=0 BASE	40	38	439	302	372	5.1
itr=250	82	70	815	578	646	6.1
itr=250 SC	354	163	1913	2174	2185	5.0
itr=0 CSB	409	258	3016	2747	2824	11.0
pkt-gen	400	360	4205			
Linux on QEMU-KVM, VALE backend, Guest-Guest						
1 VCPU	UDP8 Kpps	UDP-1460 Kpps Mbps	TCP Mbps	TCPw Mbps	TCPRR KTps	
TX = RX						
itr=0 BASE	32	31	363	559	805	7.4 *
itr=250	125	98	1143	792	859	7.1
itr=250 SC	455	186	2173	1856	2450	6.7
itr=0 CSB	1526	480	5616	4190	4206	12.3 *
2 VCPU	UDP8 Kpps	UDP-1460 Kpps Mbps	TCP Mbps	TCPw Mbps	TCPRR KTps	
TX = RX						
itr=0 BASE	77	66	767	456	548	7.3
itr=250	118	90	1059	813	900	7.0
itr=250 SC	468	286	3350	2144	3130	6.6
itr=0 CSB	1221	447	5226	4380	4849	11.2
pkt-gen	2800					

Figure 4: Performance of the various solutions with TAP or VALE as a backend, and Linux as the guest OS. Lines with an \* indicate that the receiver is not able to sustain the transmit rate.

- the **moderation delay** (**itr**);
- the use of either of the **SC** or **CSB** extensions;
- **packet size** (8 or 1460 bytes for UDP) or **write() size** for TCP (we used 1460, labeled TCP, or netperf defaults, labeled TCPw);
- the emulated device, normally **e1000** except for Figure 11 which measures the throughput with VIRTIO.

The full experimental results in a variety of configurations for Linux and FreeBSD are listed in Figures 4, 5, 10 and 11.

FreeBSD on QEMU-KVM, TAP backend, Guest-Host

1 VCPU	UDP8 Kpps	UDP-1460 Kpps	TCP Mbps	TCPw Mbps	TCPRR KTPs
TX itr=0 BASE	17	18	207	181	14.6
TX itr=250	64	55	640	267	8.9
TX itr=250 SC	237	167	1947	1060	10.1
TX itr=0 CSB	392	247	2886	1783	16.5
TX itr=250 CSB	379	244	2851	1767	7.2

2 VCPU	UDP8 Kpps	UDP-1460 Kpps	TCP Mbps	TCPw Mbps	TCPRR KTPs
TX itr=0 BASE	17	17	193	176	14.1
TX itr=250	59	53	625	251	12.2
TX itr=250 SC	244	165	1930	1068	11.6
TX itr=0 CSB	400	244	2851	1862	17.3
TX itr=250 CSB	388	243	2844	1746	7.2

FreeBSD on QEMU-KVM, TAP backend, Guest-Guest

1 VCPU	UDP8 Kpps	UDP-1460 Kpps	TCP Mbps	TCPw Mbps	TCPRR KTPs
TX = RX					
itr=0 BASE	20	19	224	185	7.1
itr=250	54	46	543	357	3.8
itr=250 SC	251	180	2109	739	3.9 *
itr=0 CSB	475	262	3055	1512	9.5 *

2 VCPU	UDP8 Kpps	UDP-1460 Kpps	TCP Mbps	TCPw Mbps	TCPRR KTPs
TX = RX					
itr=0 BASE	16	16	187	157	6.6
itr=250	50	48	564	422	3.8
itr=250 SC	250	170	1996	704	3.8 *
itr=0 CSB	464	268	3131	1350	8.0

FreeBSD on QEMU-KVM, VALE backend, Guest-Guest

1 CPU	UDP8 Kpps	UDP-1460 Kpps	TCP Mbps	TCPw Mbps	TCPRR KTPs
TX = RX					
itr=0 BASE	20	18	205	160	7.9
itr=250	91	60	697	547	4.3
itr=250 SC	459	268	3126	927	5.7 *
itr=0 CSB	838	456	5325	1835	9.9 *

2 CPU	UDP8 Kpps	UDP-1460 Kpps	TCP Mbps	TCPw Mbps	TCPRR KTPs
TX = RX					
itr=0 BASE	20	20	230	183	7.3
itr=250	85	65	764	439	4.3
itr=250 SC	440	258	3018	906	5.0 *
itr=0 CSB	817	449	5248	1684	9.1 *

pkt-gen	5000	1890	22075		
---------	------	------	-------	--	--

Figure 5: Performance of the various solutions with TAP or VALE as a backend, and FreeBSD as the guest OS. Lines with an \* indicate that the receiver is not able to sustain the transmit rate.

We will provide a detailed discussion of some of the most significant configurations in the rest of this Section.

#### 4.1 Effect of NIC improvements

We show here a few graphs to compare the effect of NIC improvements on the communication speed.

Figure 6 shows how performance is affected by various combinations of the parameters, using TAP (grey) or VALE (black) as a backend, and short UDP packets (the effect on large UDP frames is similar, as can be seen from the tables). The graph shows clearly that with a TAP backend, paravirtualization of the `e1000` NIC (column labeled `csb`) matches the performance of VIRTIO without VHOST. This was expected as the two mechanisms operate in a very similar way and have comparable overheads. As discussed, VIRTIO with VHOST has slightly superior performance because it avoids going through the I/O thread hence reduces the overall packet processing costs. However, when VALE




Figure 6: UDP throughput with various configurations.

is used as a backend, our emulated `e1000` is significantly faster than all other modes, because our modified `e1000` frontend can exploit the ability of the VALE backend to handle batches of packets in a single system call.

We note that even if we cannot use paravirtualization (which requires support in both the guest OS and to the VMM), the other options that we provide still give substantial benefits. Plain interrupt moderation, which is a VMM-only feature that has now been imported in QEMU and we can expect to be available on other hypervisors, more than doubles the packet rate. Moderation with send combining, which is a trivial device driver modification for the guest OS, almost matches the throughput of VIRTIO.

#### 4.2 Latency

Figure 7 presents the results of the “request-response” test (TCP RR) in netperf. The test, after establishing a TCP connection, exchanges one-byte packets between the peers, measuring the number of exchanges per unit of time. Hence, it measures the average latency of the path.




Figure 7: TCP RR rates with various configurations.



None of our mechanisms is designed specifically to improve latency, and in fact there is normally a tradeoff between latency and throughput. The goal of these tests is to verify that our modifications are not affecting negatively the communication delay

As expected, we see that send combining is not effective as there is normally only one packet in flight. Interrupt moderation has mixed effects, as it may slightly reduce transmit interrupts, but may also delays receive notifications. The paravirtualized mode (CSB) gives some benefits (actually halving the latency of the path) and almost matching the performance of VIRTIO. This also was expected, given the similarity of the two mechanisms; VIRTIO has a small advantage in that in the `e1000` driver the interrupt that issues the receive kick is more expensive than the one for VIRTIO (which uses MSI-X and does not need to clear a status register). As usual, VIRTIO+VHOST gets a performance boost as it avoids going the extra delay induced by the handoff to the I/O thread, which is the really expensive operation in this type of test.

Finally, we see that the VALE switch, although faster than TAP, gives only a modest improvement to the transaction rate, because the switch is responsible only for a small fraction of the overall delay of the path.

### 4.3 TCP throughput

As mentioned, TCP throughput is not the main goal of this paper. High bulk transfer rates are generally achieved by reducing the delay of the path, so that even modest window sizes suffice to keep the link streaming, and using large segments, possibly beyond the maximum frame size of the link, to reduce the CPU load on the sender and receiver; hardware support (TSO and LRO) can perform segmentation and reassembly relieving the CPU from this task, and also reducing packet rates by a factor of 10 or more.

TSO and LRO are particularly useful in the case of virtual machines residing on the same host, because the emulated device can completely skip the segmentation, and pass the entire TSO segment to the other endpoint. A similar technique can be applied to checksumming. Both optimizations




Figure 8: Guest-Guest TCP throughput with different frontend/backend (only Linux guests, no TSO).




Figure 9: TCP throughput with various guest/host combinations (only TAP backend).

are normally used by the VIRTIO driver when running on top of a TAP backend.

We are addressing TCP optimization in a separate work. In this paper, we want to verify how the proposed mechanism impact TCP throughput in absence of accelerations, hence the tests in this section are run without TSO/LRO.

Figure 8 shows the behaviour of the various solutions for TCP traffic, again between two guests, and without using TSO. As we can see, the paravirtualized e1000 almost matches VIRTIO+TAP, and is slightly less performant than VIRTIO+VHOST due to the different delay of the path (in particular VIRTIO is able to avoid computing the TCP checksum, while the e1000 device does the checksum computation in the device emulator). While increasing the TCP window/socket buffer size can mitigate the effects of delay, large windows also introduce side effects such as increased memory footprint and reduced cache locality, which at the rates we are dealing with (several Gbit/s) can still have a measurable impact on throughput.

When one of the endpoints is on the host, it can act as a faster source or sink of traffic, and put under stress the other endpoint on the guest. In Figure 9 we compare the throughput of a connection when the sender and the receiver are placed on the host or on the guest. What is evident here is that the basic e1000 device is unable to send at high speed (columns Guest→Host and Guest→Guest) but the receive direction has much better performance (column Host→Guest), presumably because packets arrive in bursts. As we bring in interrupt moderation and send combining, the transmit direction starts exhibiting good performance, and the paravirtualized mode performs almost as good as VIRTIO.

Similarly, the receive direction exhibits a very good receive performance when the transmitter is on the host (column Host→Guest), especially in paravirtualized mode.

### 4.4 VALE backend

We introduced the VALE switch to replace the TAP backend because the latter had become a bottleneck in the communication between the virtual machines.

The throughput between two ports of the VALE virtual switch, on our test machine, is approximately 20 Mpps with small UDP frames, and 70 Gbit/s with 1500-byte frames. When used to interconnect VMs, the VALE port is attached to the VMM backend and from there to the frontend. These components introduce extra data copies and delays in the communication, leading to a reduction in the throughput<sup>1</sup>.

An estimate of the maximum guest-to-guest throughput can be determined by running instances of netmap-based traffic sources and sinks (**pkt-gen**) within the two guest. This experiment achieved about 5 Mpps with 64-byte frames, and 1.89 Mpps with 1500-byte frames, corresponding to 22.7 Gbit/s.

All these numbers are significantly higher than the packet rates we experience using regular sockets. This suggests that the VALE switch is not acting as a bottleneck for the system. Indeed, as shown in Figure 6 (the black column for the CSB case) even socket based applications running in the two guests can exchange about 1.2 Mpps, which is the same rate we achieve for the same system on bare metal.

## 4.5 Summary

In this paragraph we summarize the proposed optimizations and the situations to which they apply. Recall that this work is mostly focused on improving packet rate.

- When only the hypervisor can be modified, implementing an interrupt moderation mechanism improves the TX and RX packet rates by amortizing the interrupt overhead over many processed packets. It is not even necessary that the real device or guest OS actually supports a moderation mechanisms.

Effect: up to two-fold speed-up on TX side, and up to 10x improvement of the maximum RX rate.

- When only the guest driver can be modified, send combining (SC) can improve the TX rate, but only when the guest manages to enqueue many TX requests to the device before a TX completion interrupts comes. The presence of multiple VCPUs, and especially the availability of interrupt moderation boost the effectiveness of this mechanism.

Effect: up to 5-10x TX speedup.

Alternatively, lazy TX completion can be implemented, to suppress most of the TX interrupts by keeping them disabled as far as the device TX queue(s) are not full. Benefits are more limited (up to 2x TX speedup) and the mechanism is not compatible with send combining.

Pure polling is another solution that completely suppress device interrupts, dramatically improving TX/RX performance. However, it may affects receive latency.

Effect: over 10x speedup on TX side, and up to 20x increase of the maximum RX rate.

- When both the driver and the hypervisor can be modified, a simple paravirtualized extension can remove most

<sup>1</sup>the delay has an impact too, because the ports of the switch have a bounded queue and new packets cannot be sent until the queue has been drained. The mechanism is similar to what happens for window-limited TCP connections.

of the virtualization overheads, and reach the highest possible packet-rates, without negatively affecting latency.

Effect: over 10x speedup on TX side, and up to 30x increase of the maximum RX rate.

In addition to that, when the hypervisor can be modified, adding a backend support for VALE usually improves performances, being generally faster than the popular TAP backend. In our experiments, the speedups for the TX side become 4x with interrupt moderation, 14x for send combining, and 40x with a paravirtualized device.

Linux on VirtualBox and VMware, 2 CPU

			UDP8	UDP-1460	TCP	TCPw	TCPRR
			Kpps	Kpps	Mbps	Mbps	Ktps
VBBox	TX	GH	22	23	264	84	633
VBBox	TX	GG	22	21	244	1121	1255
VMware	TX	GH	52	51	590	250	1332
VMware	TX	GG	65	64	748	3375	4138

FreeBSD on VirtualBox, vboxnet, Guest-Host

1 cpu, std	24	24	273	219	666	16.9
1 cpu, SC	24	24	273	232	928	17.0
2 cpu, std	60	58	676	570	690	14.7
2 cpu, SC	225	176	2064	1060	1100	14.7

Backend performance:

pktgen, tx	540	470.0	5500
pktgen, rx	670	270.0	3153

Figure 10: Performance of other VMMs (e1000 device). SC can help even without VMM modifications.

## 4.6 Comparison with other solutions

The huge performance improvements that we see with respect to the baseline could be attributed to inferior performance of QEMU-KVM compared to other solutions, but this is not the case.

Figure 10 shows that the performance of VirtualBox and VMware Player in a similar configuration is comparable with our BASE configuration with QEMU. We hope to be able, in the future, to run comparisons against other VMMs.

We also tried to use Send Combining on an unmodified VirtualBox (without moderation), and at least in the 2-VCPU case, see Figure 10, this is still able to give significant performance improvements, presumably due to a reduction in the number of VM exits on transmissions.

## 5. CONCLUSIONS AND FUTURE WORK

We have shown how some small and simple modifications to hypervisors and device drivers can go a long way into making network performance on virtual machine very close to that of bare metal, in particular for the case of high packet rate workloads.

Our extensions to the e1000 (and r8169) drivers can already improve the network performance of guest OSes by a large factor. The use of VALE as a virtual switch permits a much faster interconnection between virtual machines and the host. Our modifications, which are publicly available, will hopefully contribute to deploy high performance SDN components in virtualized environments, and also help the

#### Linux on QEMU, VIRTIO, no accelerations

			UDP8 Kpps	UDP1460 Kpps	Mbps	TCP Mbps	TCPw Mbps	TCPRR Ktps
1 CPU	GH	tap	342	320	3742	3490	3485	24.3
2 CPU	GH	tap	334	312	3653	3346	3335	26.7
1 CPU	GG	tap	424	394	4610	3475	3480	14.5
2 CPU	GG	tap	417	367	4294	3359	3416	13.8
1 CPU	GG	vale	981	768	9184	4526	4532	15.3 *
2 CPU	GG	vale	875	774	9044	4346	3682	14.6 *

#### Linux on QEMU, VIRTIO, VHOST

			UDP8 Kpps	UDP1460 Kpps	Mbps	TCP Mbps	TCPw Mbps	TCPRR Ktps
1 CPU	GH		644	564	6594	7144	7303	37.0
2 CPU	GH		642	576	6724	7039	6880	35.8
1 CPU	GG		528	477	5572	5893	5868	22.7
2 CPU	GG		871	768	8976	4450	4539	21.8 *

#### Linux on QEMU, VIRTIO, TSO

			UDP8 Kpps	UDP1460 Kpps	Mbps	TCP Mbps	TCPw Mbps	TCPRR Ktps
1 CPU	GH	tap				11329	21478	
1 CPU	GH	vhost				13107	26750	
2 CPU	GH	tap				10809	21293	
2 CPU	GH	vhost				12687	26519	
1 CPU	GG	tap				19510	19435	13.7
1 CPU	GG	vhost				17286	33015	22.6
2 CPU	GG	tap				16558	18930	13.1
2 CPU	GG	vhost				12391	26867	21.6

Figure 11: Performance with VIRTIO and TSO.

study and development of high speed protocol and applications over virtual machines.

At the time of this writing, the interrupt moderation emulation has been already imported in the QEMU distribution, and we are working to integrate the other features in FreeBSD, Linux and QEMU.

While our performance numbers are excellent for UDP, the fact that we do not yet exploit the guest TSO support to obtain a large effective MTU (see Section 4) is severely penalizing for the case of bulk TCP connections. Figure 11 shows that TSO increments the TCP throughput up to 3 times for VIRTIO. We are extending the e1000 emulator and the VALE switch so that they can support large segments and scatter/gather I/O. Preliminary measurements indicate that we can match the VIRTIO results even in this case.

We are also working on reducing the latency of the path between virtual machines, so that TCP communications do not require exceedingly large window sizes to achieve decent throughput.

## Acknowledgements

This work has been supported by the European Commission and developed in the contexts of projects CHANGE (INFO-ICT-257422) and OPENLAB (INFO-ICT-287581).

## 6. REFERENCES

- [1] AGESEN, O., GARTHWAITE, A., SHELDON, J., AND SUBRAHMANYAM, P. The evolution of an x86 virtual machine monitor. *SIGOPS Oper. Syst. Rev.* 44, 4 (Dec. 2010), 3–18.
- [2] AGESEN, O., MATTSO, J., RUGINA, R., AND SHELDON, J. Software techniques for avoiding

hardware virtualization exits. *USENIX ATC'12*, USENIX Association, pp. 35–35.

- [3] AMD. Secure virtual machine architecture reference manual. <http://www.mimuw.edu.pl/~vincent/lecture6/sources/amd-pacifica-specification.pdf>, 2005.
- [4] BEN-YEHUDA, M., MASON, J., KRIEGER, O., XENIDIS, J., DOORN, L. V., MALLICK, A., AND WAHLIG, E. Utilizing IOMMUs for Virtualization in Linux and Xen. In *Proceedings of the Linux Symposium* (2006).
- [5] BUGNION, E., DEVINE, S., ROSENBLUM, M., SUGERMAN, J., AND WANG, E. Y. Bringing virtualization to the x86 architecture with the original vmware workstation. *ACM Trans. Comput. Syst.* 30, 4 (Nov. 2012), 12:1–12:51.
- [6] CHISNALL, D. *The Definitive Guide to the Xen Hypervisor*. Prentice Hall, 2007.
- [7] DERI, L. PFRING DNA page. [http://www.ntop.org/products/pf\\_ring/dna/](http://www.ntop.org/products/pf_ring/dna/).
- [8] GORDON, A., AMIT, N., HAR'EL, N., BEN-YEHUDA, M., LANDAU, A., SCHUSTER, A., AND TSAFRIR, D. ELI: bare-metal performance for I/O virtualization. *SIGARCH Comp. Arch. News* 40, 1 (Mar. 2012), 411–422.
- [9] GORDON, A., HAR'EL, N., LANDAU, A., BEN-YEHUDA, M., AND TRAEGER, A. Towards exitless and efficient paravirtual i/o. *SYSTOR '12*, ACM, pp. 4:1–4:6.
- [10] GUAN, H., DONG, Y., MA, R., XU, D., ZHANG, Y., AND LI, J. Performance Enhancement for Network I/O Virtualization with Efficient Interrupt Coalescing and Virtual Receive Side Scaling. *IEEE Transactions on Parallel and Distributed Systems*.
- [11] HAJNOCZI, S. QEMU Internals: vhost architecture. <http://blog.vmsplce.net/2011/09/qemu-internals-vhost-architecture.html>, 2011.
- [12] HAN, S., JANG, K., PARK, K., AND MOON, S. Packetshader: a gpu-accelerated software router. *ACM SIGCOMM Computer Communication Review* 40, 4 (2010), 195–206.
- [13] INTEL. Intel virtualization technology. *Intel Tech. Journal* 10 (Aug. 2006).
- [14] INTEL. Intel data plane development kit. <http://edc.intel.com/Link.aspx?id=5378> (2012).
- [15] MOGUL, J., AND RAMAKRISHNAN, K. Eliminating receive livelock in an interrupt-driven kernel. *ACM Transactions on Computer Systems (TOCS)* 15, 3 (1997), 217–252.
- [16] RIZZO, L. Netmap home page. *Università di Pisa*, <http://info.iet.unipi.it/~luigi/netmap/>.
- [17] RIZZO, L. Polling versus interrupts in network device drivers. *BSDConEurope 2001* (2001).
- [18] RIZZO, L. netmap: A Novel Framework for Fast Packet I/O. In *USENIX ATC'12* (2012), Boston, MA, USENIX Association.
- [19] RIZZO, L., AND LETTIERI, G. The VALE Virtual Local Ethernet home page. <http://info.iet.unipi.it/~luigi/vale/>.

- [20] RIZZO, L., AND LETTIERI, G. VALE, a switched ethernet for virtual machines. In *Proceedings of the 8th international conference on Emerging networking experiments and technologies* (New York, NY, USA, 2012), CoNEXT '12, ACM, pp. 61–72.
- [21] RUSSELL, R. virtio: towards a de-facto standard for virtual I/O devices. *SIGOPS Oper. Syst. Rev.* 42, 5 (July 2008), 95–103.
- [22] SALIM, J. H., OLSSON, R., AND KUZNETSOV, A. Beyond softnet. In *Proceedings of the 5th annual Linux Showcase & Conference - Volume 5* (Berkeley, CA, USA, 2001), ALS '01, USENIX Association, pp. 18–18.
- [23] SCHULTZ, M., AND CROWLEY, P. Performance analysis of packet capture methods in a 10 gbps virtualized environment. In *Computer Communications and Networks (ICCCN), 2012 21st International Conference on* (30 2012-aug. 2 2012), pp. 1–8.
- [24] SOLARFLARE. Openonload. <http://www.openonload.org/> (2008).
- [25] SUGERMAN, J., VENKITACHALAM, G., AND LIM, B.-H. Virtualizing I/O Devices on VMware Workstation's Hosted Virtual Machine Monitor0. In *USENIX ATC 2002* (Berkeley, CA, USA, 2001), USENIX Association, pp. 1–14.
- [26] VMWARE. Performance Evaluation of VMXNET3. [http://www.vmware.com/pdf/vsp\\_4\\_vmxnet3\\_perf.pdf](http://www.vmware.com/pdf/vsp_4_vmxnet3_perf.pdf).